

Recurrent Neural Networks

Madhumitha Sakthi, Dr. Edison Thomaz - EE380L Data Mining- Spring 2020

Topics covered

- ▶ Sequence Modeling
- ▶ RNN - Introduction
- ▶ Traditional RNN Architecture
- ▶ Types of RNN and it's usage
- ▶ RNN Variants
- ▶ BPTT
- ▶ RNN Limitations
- ▶ LSTM
- ▶ GRU
- ▶ Applications - LSTM, GRU
- ▶ Seq2Seq
- ▶ Attention Networks
- ▶ Introduction to ^^NLP^^

Sequence Modeling

- ▶ It is the task of predicting the next timestep based on the past
 - ▶ Text Autocomplete predicts the word based on letters
- ▶ $x_t \sim p(x_t | x_{t-1}, \dots, x_1)$ i.e., predict the current value based on a distribution from the past
- ▶ What if we have (t-1) observations that makes it computationally expensive?
 - ▶ 1st Strategy: Assume that you need only 'm' observations from the past: makes 'm' constant and train autoregressive models
 - ▶ 2nd Strategy: Learn a summary about the past and keep updating this: $x_t = p(x_t | x_{t-1}, h_t)$ Latent Autoregressive models

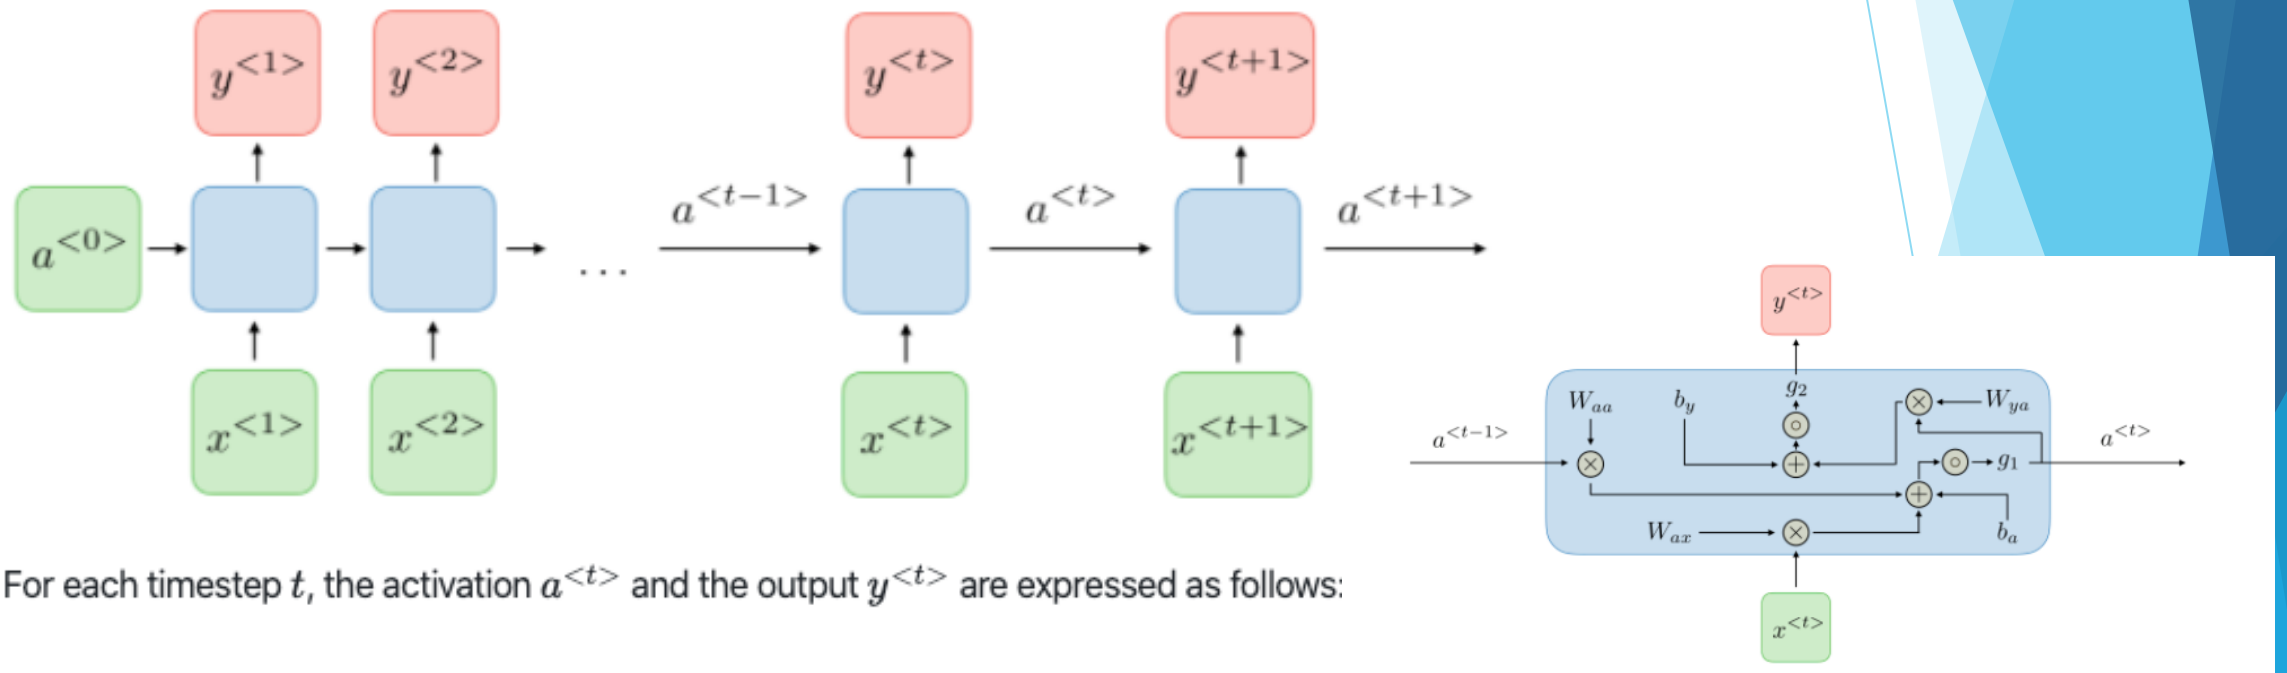
RNN - Introduction

- ▶ A simple neural network for sequence modeling
- ▶ The hidden state of each step would depend on the previous step
- ▶ The hidden state h_t acts as a memory about the past
- ▶ There are various types of RNN widely used in various NLP applications
- ▶ Variants in the internal architecture comes with added computational costs

- ▶ Unlike the regular GD, we would use BPTT (BackPropagation Through Time) to update the RNN weights

Traditional RNN Architecture


Previous outputs to be used as inputs while having hidden states



$$a^{<t>} = g_1(W_{aa}a^{<t-1>} + W_{ax}x^{<t>} + b_a) \quad \text{and} \quad y^{<t>} = g_2(W_{ya}a^{<t>} + b_y)$$


Types of RNN and it's usage

One-to-many
 $T_x = 1, T_y > 1$




Music generation

Many-to-many
 $T_x = T_y$




Name entity recognition

One-to-one
 $T_x = T_y = 1$




Traditional neural network

Many-to-one
 $T_x > 1, T_y = 1$



Sentiment classification

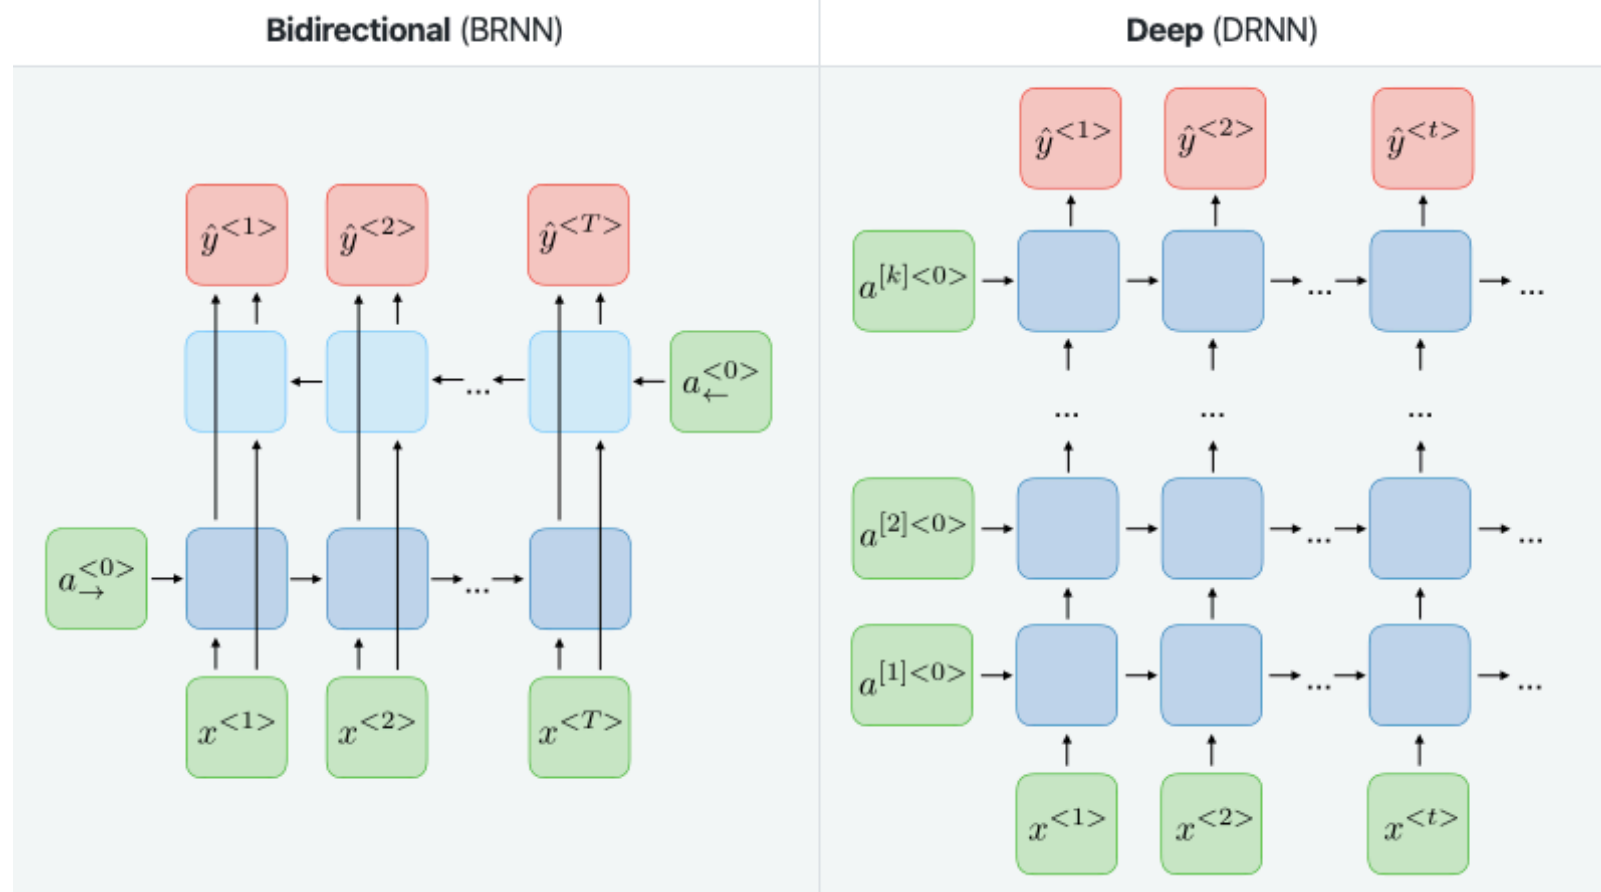
Many-to-many
 $T_x \neq T_y$



Machine translation

RNN Variants

Two commonly used BRNN variants



BPTT

- ▶ The hidden state depends on previous hidden states
- ▶ Cannot apply regular backpropagation on this

$$a^{<t>} = g_1(W_{aa}a^{<t-1>} + W_{ax}x^{<t>} + b_a) \quad \text{and} \quad y^{<t>} = g_2(W_{ya}a^{<t>} + b_y)$$


$$\mathcal{L}(\hat{y}, y) = \sum_{t=1}^{T_y} \mathcal{L}(\hat{y}^{<t>}, y^{<t>})$$

$$\frac{\partial \mathcal{L}^{(T)}}{\partial W} = \sum_{t=1}^T \frac{\partial \mathcal{L}^{(T)}}{\partial W} \Big|_{(t)}$$

- ▶ Unroll in time and compute the derivative at each timestep and sum them

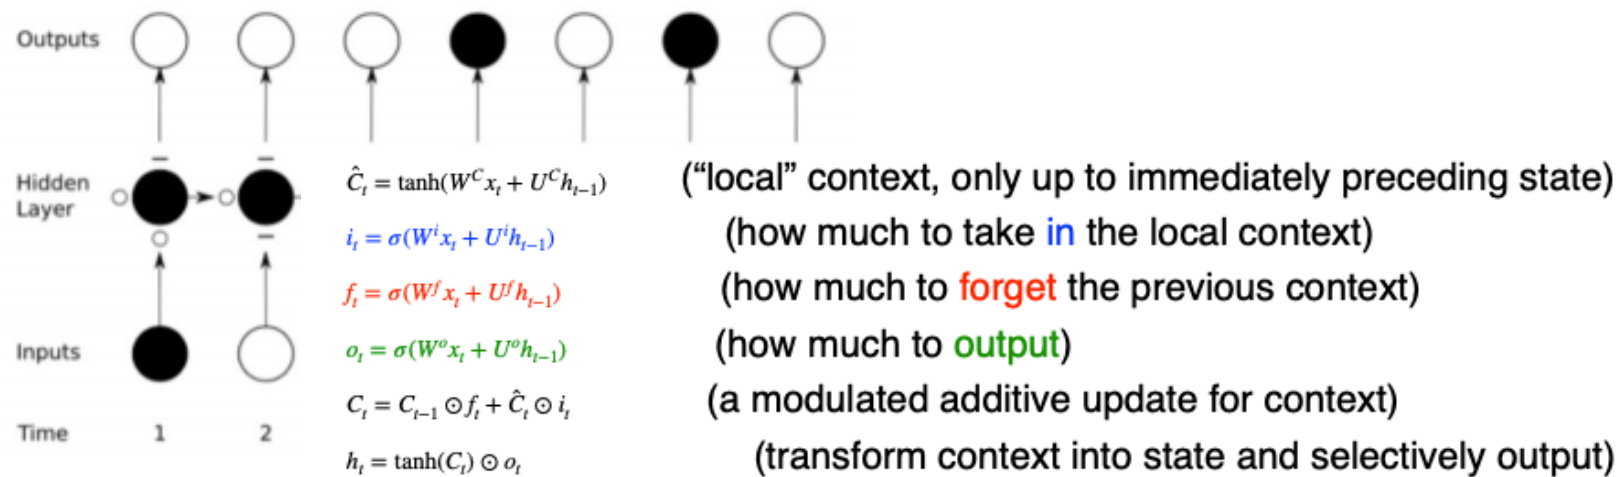
RNN Limitations

- ▶ Sensitivity of hidden states and outputs on a given input becomes weaker as we move away from it along the sequence (weak memory)
- ▶ New inputs “overwrite” the activations of previous hidden states
- ▶ Repeated multiplications can cause the gradients to vanish or explode



LSTM

- ▶ Captures long-range dependencies
- ▶ Open gate by 'o' and closed gate by '-'



- ▶ State updated are now additive and not multiplicative
- ▶ State updated were multiplicative in RNNs

GRU

- ▶ The update gate
- ▶ Then the reset gate
- ▶ There is no longer a separate cell state

$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$

$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + Ux_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

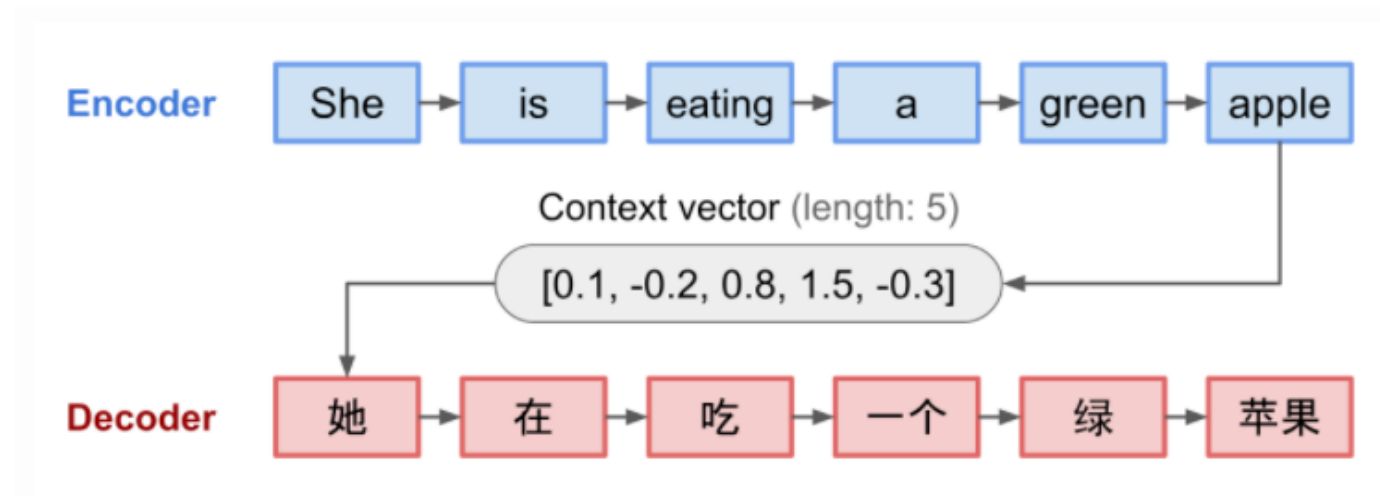
- ▶ GRU vs LSTM: GRUs are easier to train and less complex than LSTM

Applications - LSTM, GRU

- ▶ LSTM, GRU are widely used in various NLP Application
- ▶ Bi-LSTMs are used for Speech Recognition
- ▶ Machine translation - Translate text or speech from one language to another
- ▶ They are used for stock price prediction
- ▶ RNNs are used in BCI applications - Attention detection from EEG

Seq2Seq Models

- ▶ Encode-Decoder Model to predict a sequence from a sequence - Machine Translation



- ▶ This fixed-length context vector cannot remember long sequences

Attention Networks

- ▶ The alignment score is parametrized by a feed-forward network
- ▶ This network would be jointly trained with the other parameters

$$\mathbf{c}_t = \sum_{i=1}^n \alpha_{t,i} \mathbf{h}_i$$

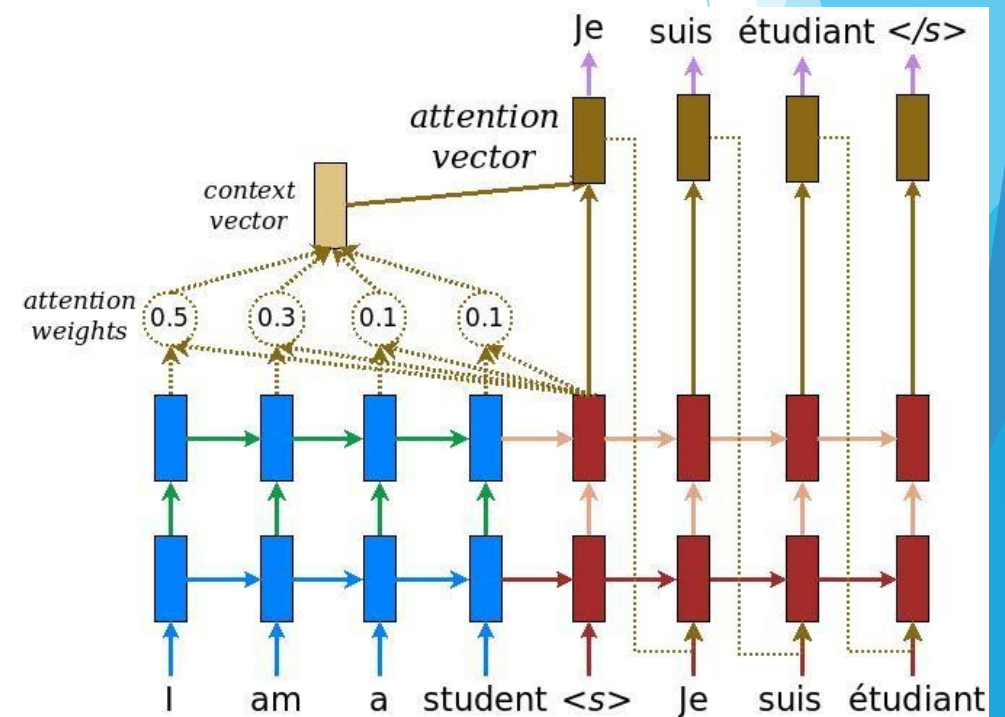
; Context vector for output y_t

$$\alpha_{t,i} = \text{align}(y_t, x_i)$$

; How well two words y_t and x_i are aligned.

$$= \frac{\exp(\text{score}(s_{t-1}, \mathbf{h}_i))}{\sum_{i'=1}^n \exp(\text{score}(s_{t-1}, \mathbf{h}_{i'}))}$$

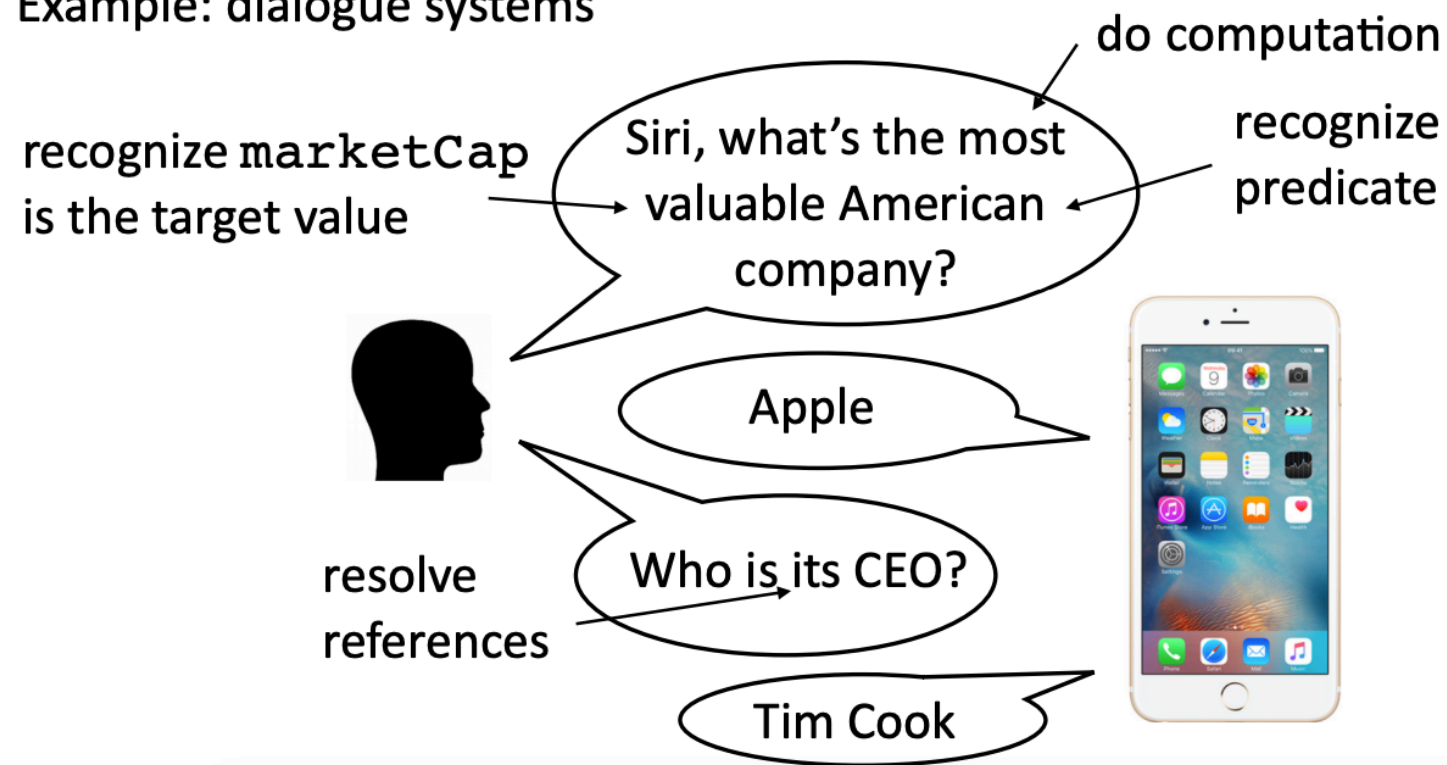
; Softmax of some predefined alignment score..



Introduction to ^^NLP^^

Be able to solve problems that require deep understanding of text

Example: dialogue systems

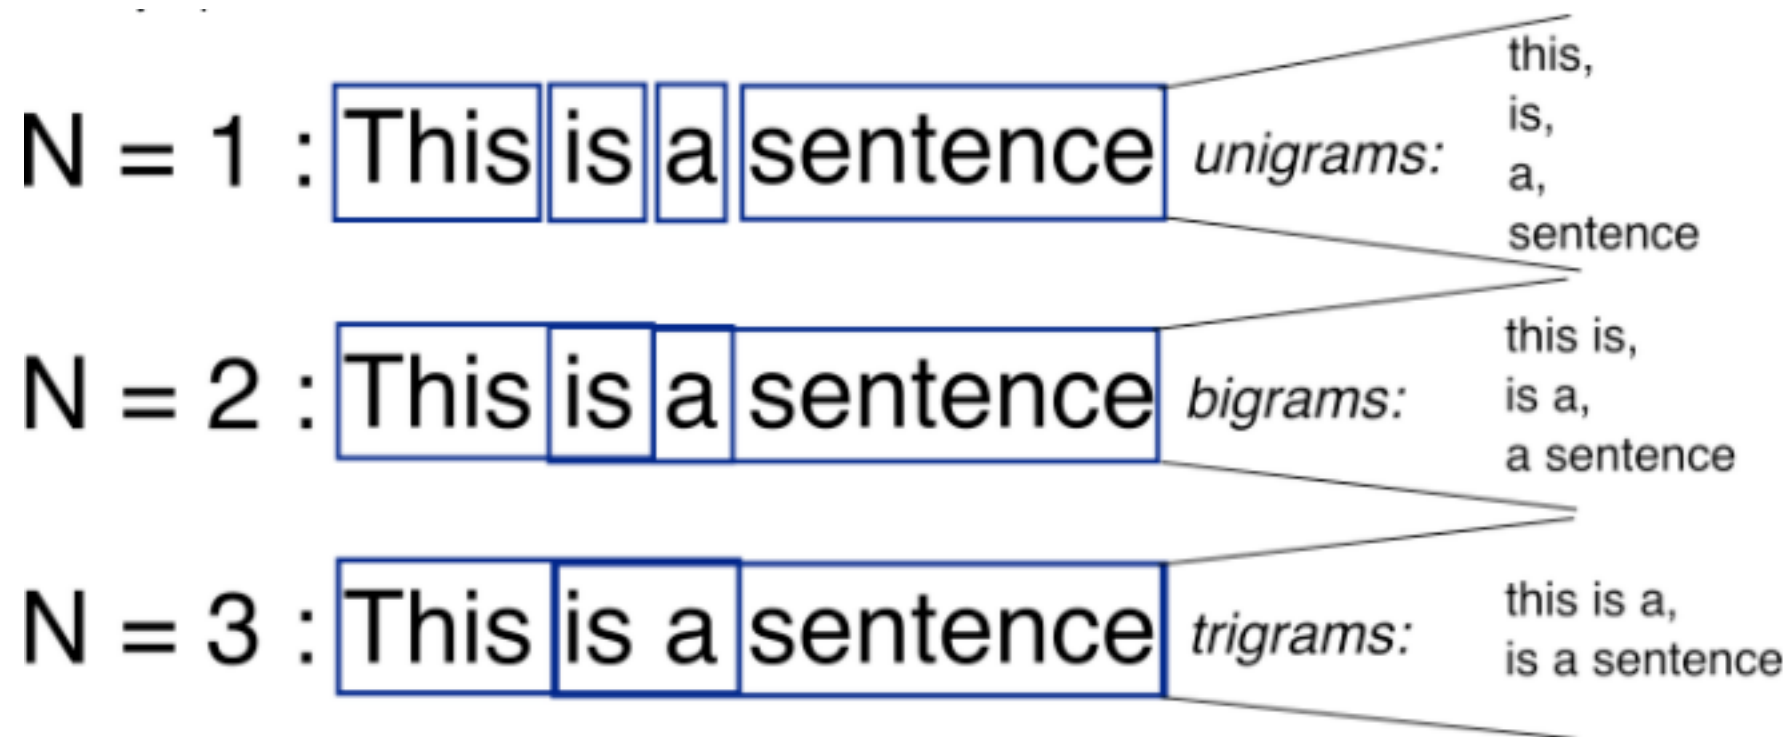


Sentiment Analysis on Text - Example

- ▶ Predict if a given sentence is positive, negative or neutral
- ▶ Text Preprocessing steps:
 - ▶ Text normalization: lower case all letters, punctuation removal, abbreviations, remove stop words, sparse words like a, an, he, is
 - ▶ Tokenization involves splitting words in a sentence as individual words
 - ▶ POS Tags: Parts of Speech tags as nouns, verbs, adjective and others
 - ▶ Stemming: Convert sentences to their root word
 - ▶ There are several types of stemming algorithms
 - ▶ There are sever type of stem algorithm

Sentiment Analysis on Text - Example

► N-gram features:



Sentiment Analysis on Text - Example

- ▶ Word2Vec
 - ▶ Statistical method to learn individual vectors for each word
 - ▶ The Continuous Bag-of-Words model learns the vector by predicting the current word based on its context
- ▶ Embedding layer:
 - ▶ This layer is learned jointly with the other layers
 - ▶ The size of the vector space is set to 50 or 100 or 300 dimensions
 - ▶ The input would be given as one-hot vector encoded words
 - ▶ Requires lot of training data but, can be trained to the specific NLP task data